

IMPERIAL

# Hand-writing Recognition With IMU and Force Sensors

## NeverLatex Final Presentation

Antonio Tarizzo

Tuna Kisaaga

Fajar Kenichi Kusumah Putra

21/03/2025

# Team



**Antonio Tarizzo**



**Tuna Kısaağa**



**Fajar Kenichi  
Kusumah Putra**

# Agenda

- 1 Introduction
- 2 Literature Review
- 3 Hardware Design
- 4 Methodology
- 5 Experimental Results
- 6 Workplan
- 7 Conclusion
- 8 Video
- 9 References

# Introduction

Scope and Objectives, Achievements,  
Background and Context

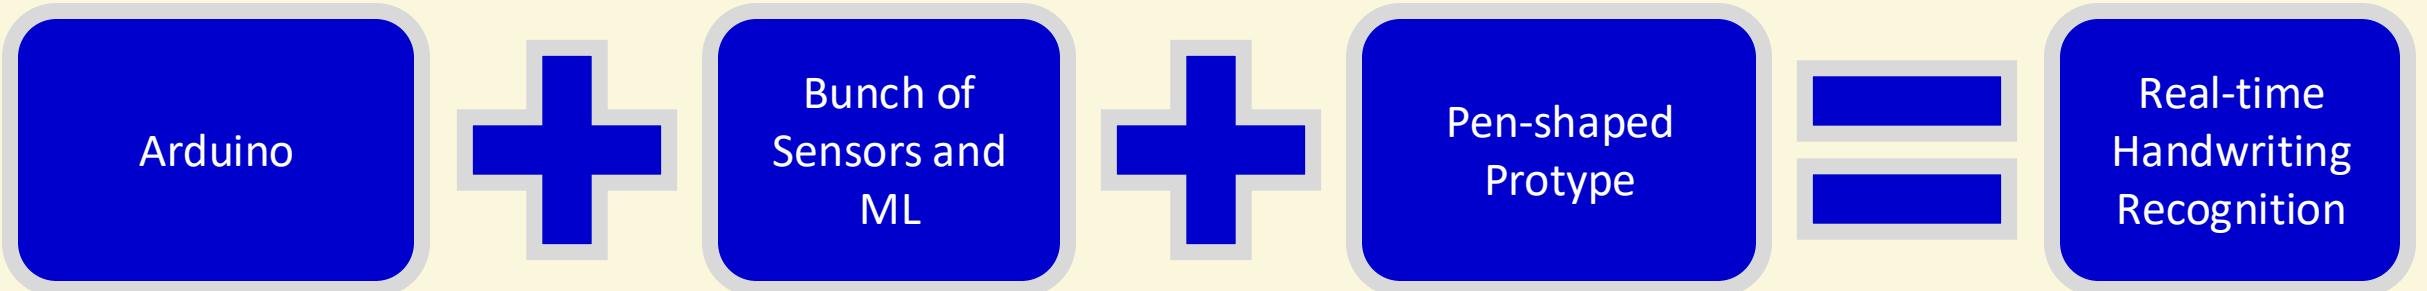
# Background and Context

## In General

This project proposes developing a **sensor-equipped handwriting recognition system** that **converts handwritten text into LaTeX format in real-time**.

Using Arduino, a pen with IMU (Inertial Measurement Units), and force strip sensors, we aim to provide a low-cost solution that eliminates the time spent to transfer information from paper to digital documents.

The project leverages machine learning to ensure high accuracy in detecting handwritten digits, text, equations and more.



# Scope, Objectives and Achievements

## Scope

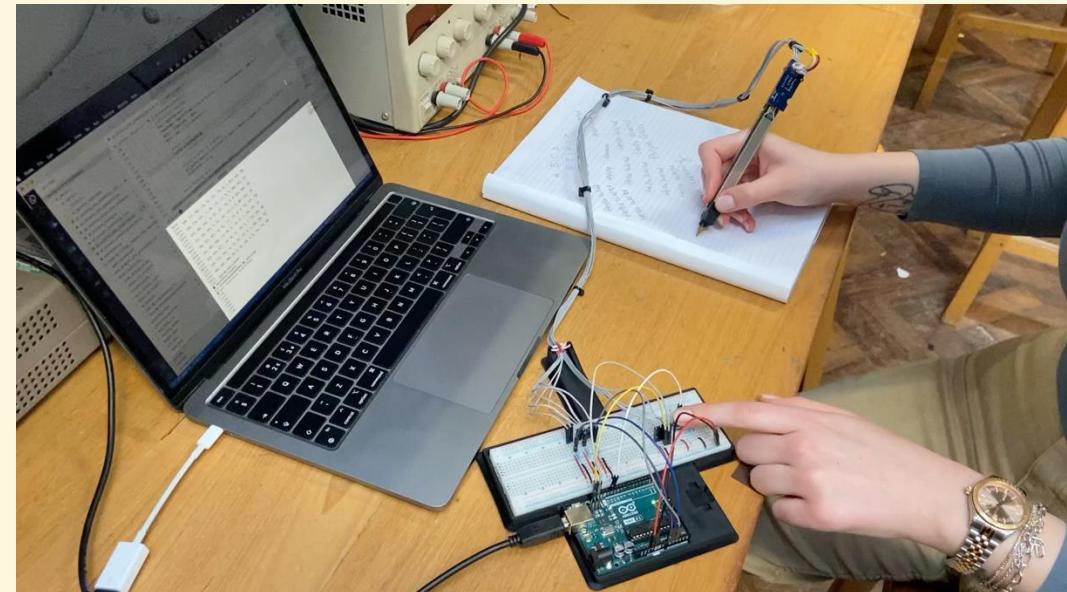
Develop a handwriting recognition system that converts handwritten text (letters, digits) and mathematical notation into LaTeX.

## Objectives

### Objectives

- 1. Handwriting Recognition:** Accurately identify digits, letters (including case), and mathematical expressions.
- 2. LaTeX Document Generation:** Automatically produce formatted LaTeX documents with sections, subsections, bulleted lists, and math environments.
- 3. Adaptive Learning & Performance:** Improve recognition accuracy and speed over time, aiming for a 10% accuracy boost and 30% speed increase in Phase 2.
- 4. Erasure & Correction:** Detect and reflect erasures in real time, with potential use of LLMs for post-recognition correction.

## Achievements



**Accuracy Achieved: 68.78%**

# Literature Review

## Reviewed Papers and Sources

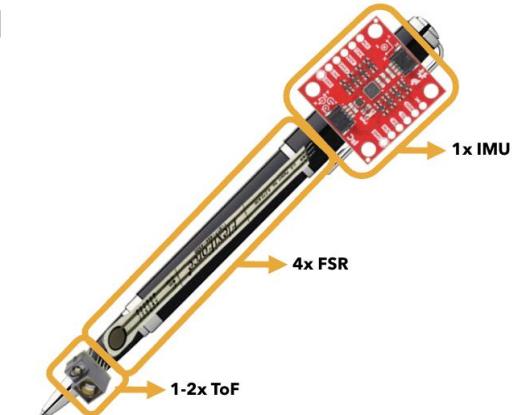
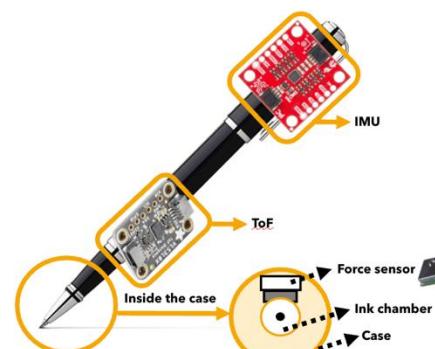
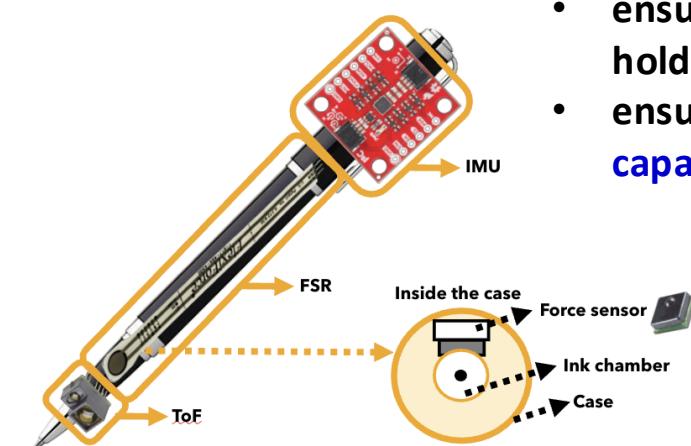
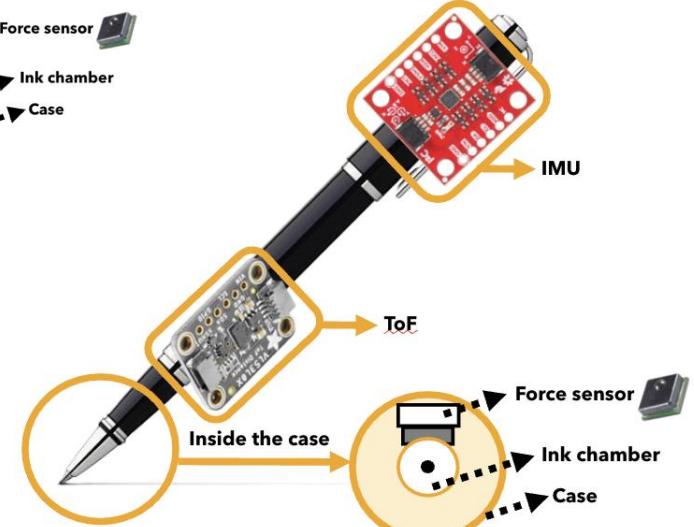
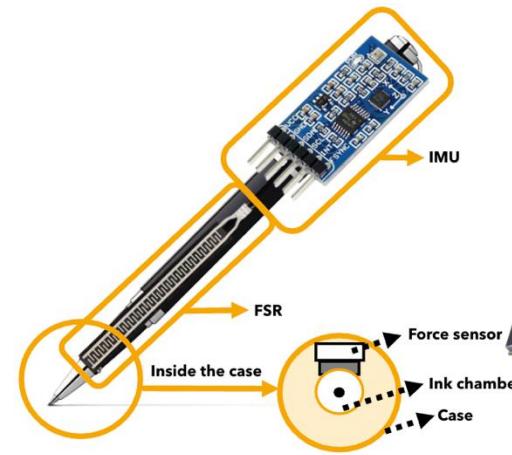
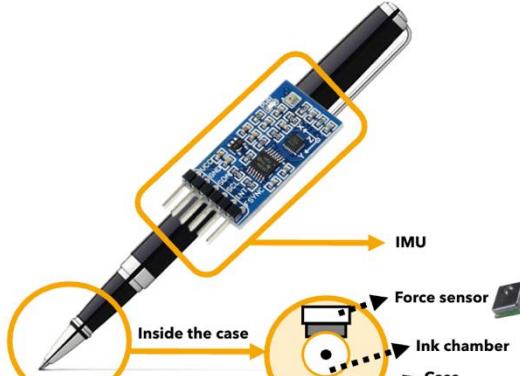
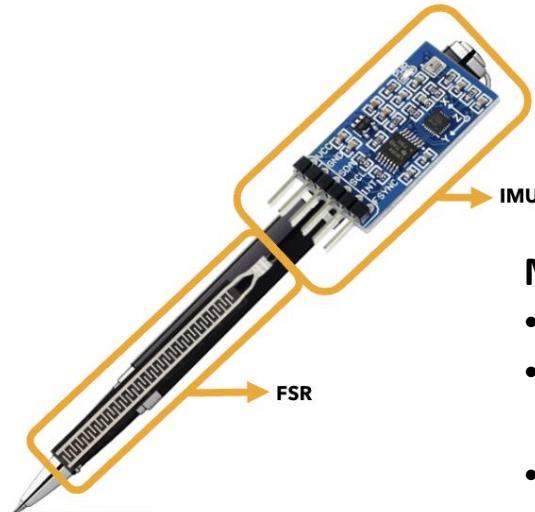
# Literature Review Key Findings

Things To Note	Action Based on Findings
<p><b>Input Data:</b></p> <ul style="list-style-type: none"><li>• Spatial data from IMU Sensors [6]</li><li>• Force data from Force Sensors [7]</li></ul> <p><b>Model Types:</b></p> <ul style="list-style-type: none"><li>• CNN [5]</li><li>• LSTM [7]</li><li>• CLDNN [5]</li></ul> <p><b>Loss Functions:</b></p> <ul style="list-style-type: none"><li>• Categorical Cross Entropy [3]</li><li>• Connectionist Temporal Classification (CTC) [5]</li></ul>	<ol style="list-style-type: none"><li>1. Implement a CNN and CLDNN based model on ONHW dataset.</li><li>2. Finetune the model on that dataset.</li><li>3. Create initial hardware prototype.</li><li>4. Collect data using the prototype.</li><li>5. Overall test (model, prototype, data collected).</li><li>6. Crosscheck results.</li><li>7. Repeat the process using the final product.</li></ol>

# Hardware Design

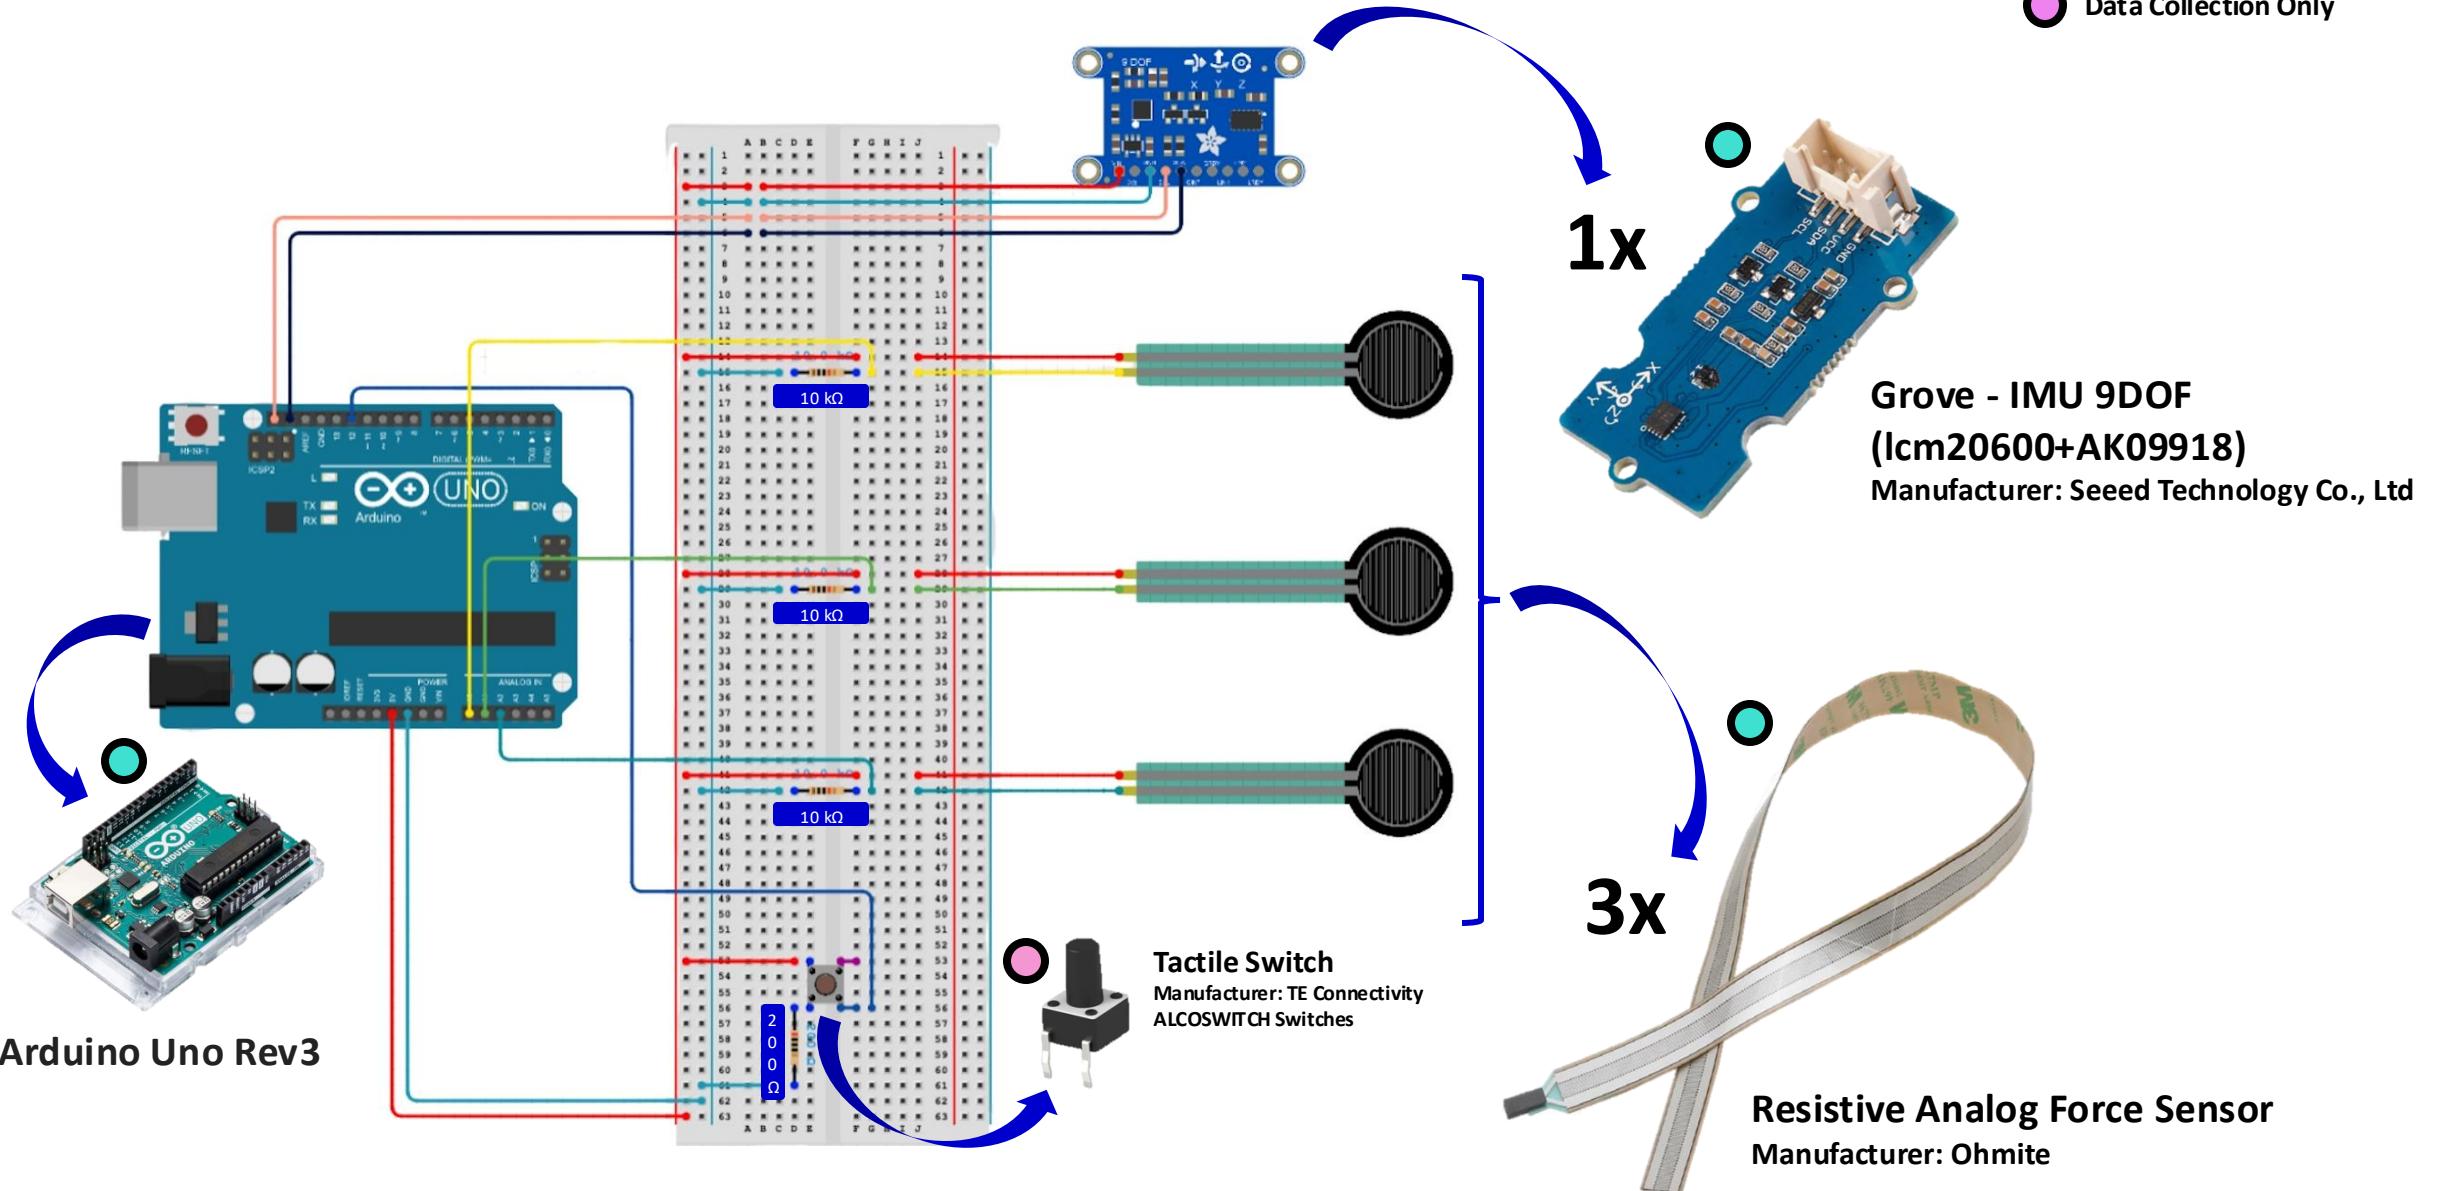
## Design Choices, Challenges and Constraints

# Considerations of Hardware Design



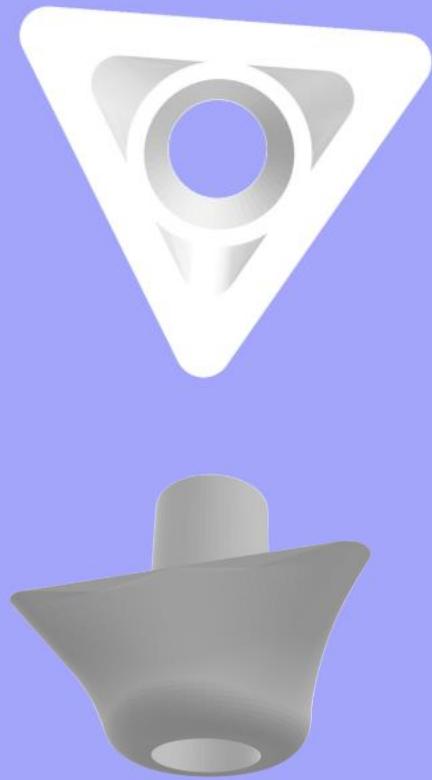
# Electronic Design Choices

- Data Collection, Training & Real-time Prediction
- Data Collection Only

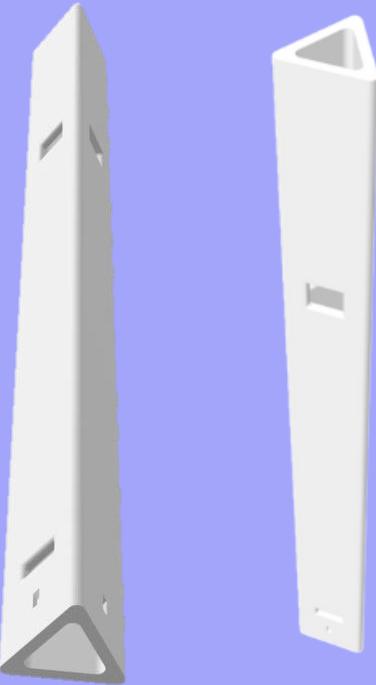


# 3D Printing Design

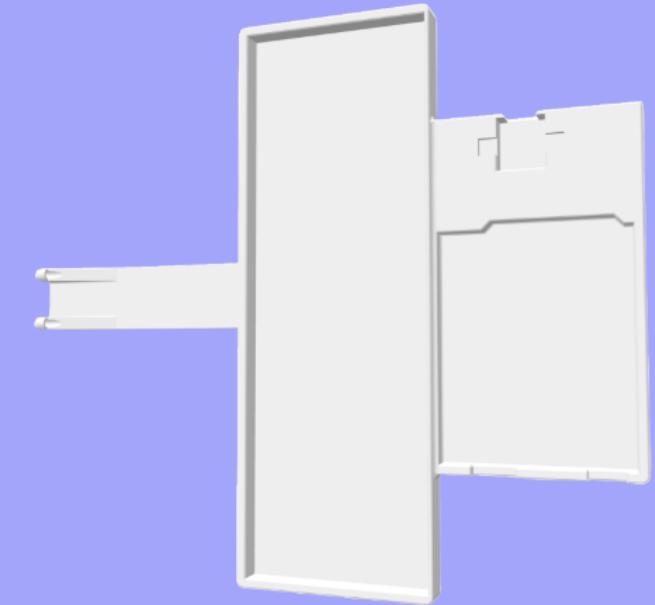
Pen Tip (for the Cartridge)



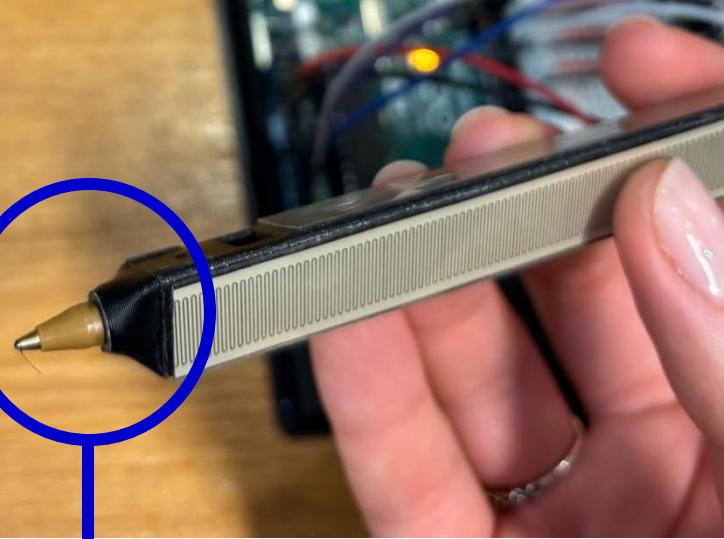
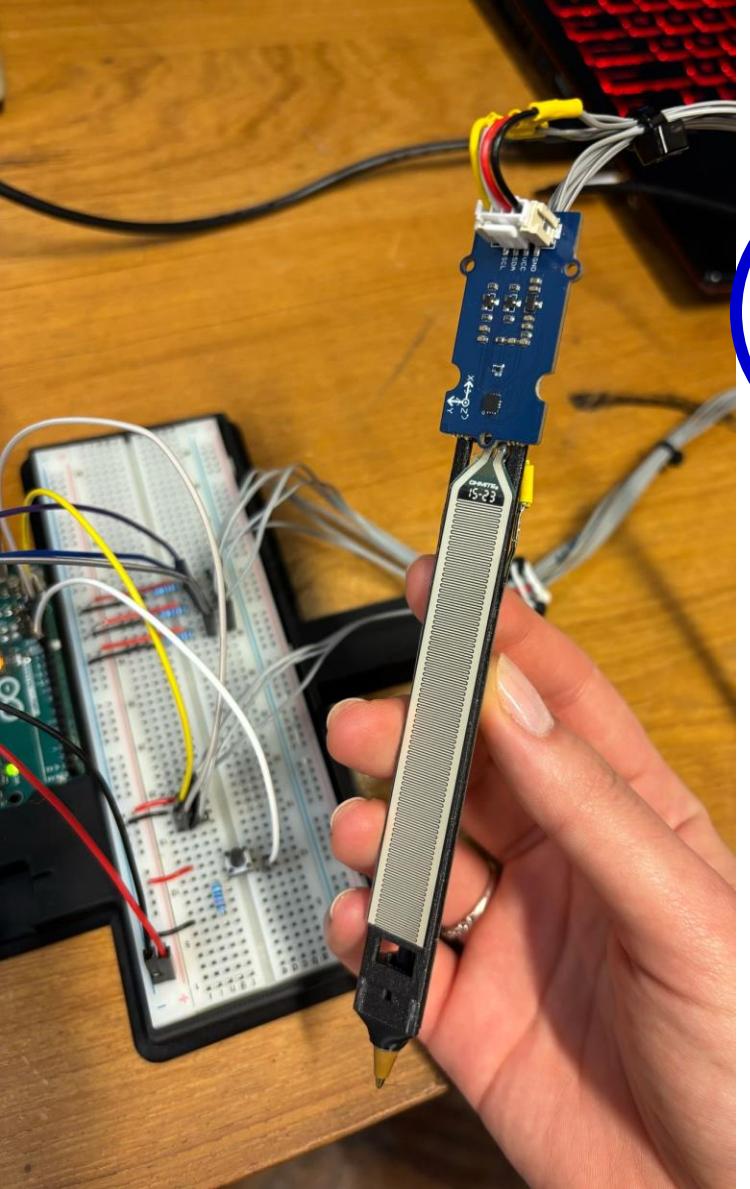
Pen Body



Arduino and Breadboard Holder

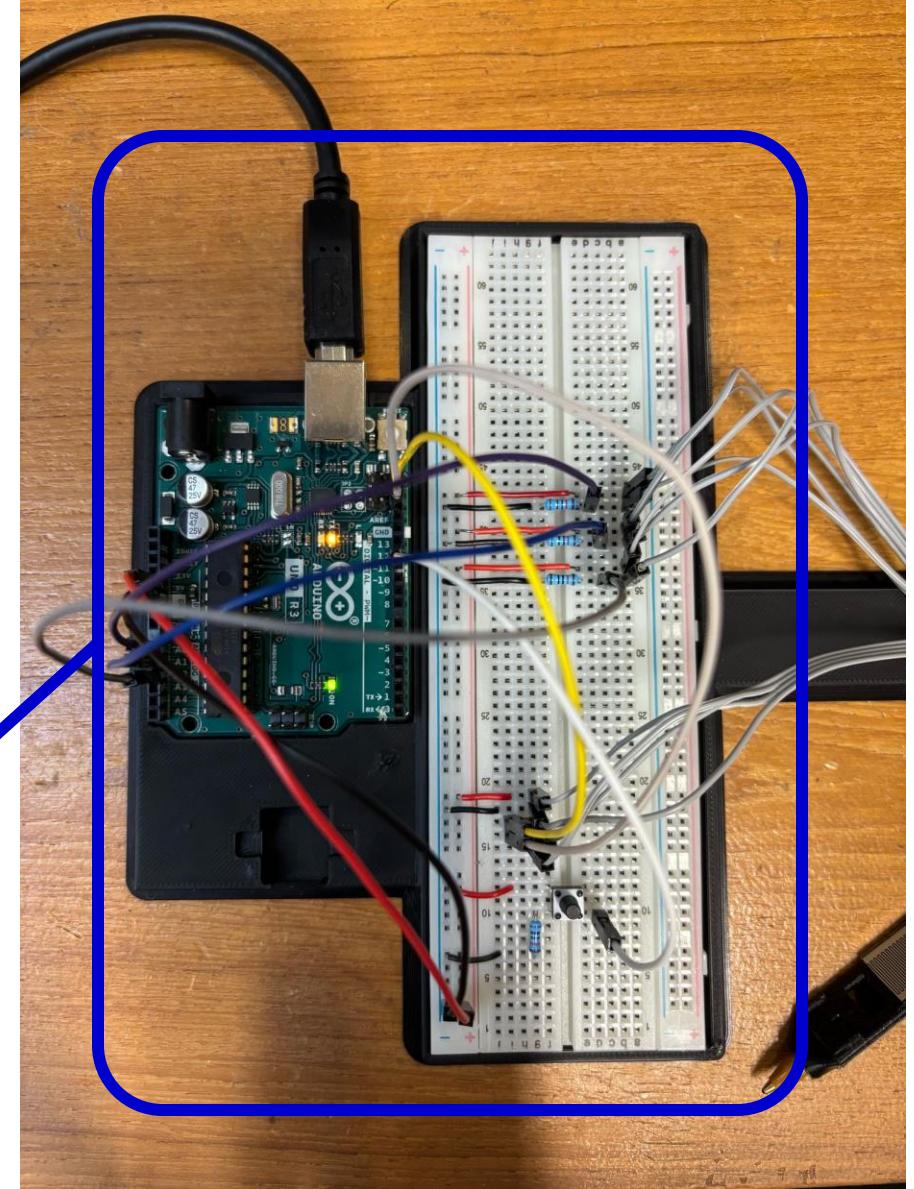


# Final Design



Pen cartridge

Arduino and  
breadboard  
holder to  
immobilise the  
writing setup



# Methodology

## Data Collection, Preprocessing, Models and Methods, Real-time Prediction

## Data Collection



## **Target characters, included in data collection**



## Only data collection

# English Alphabet

## Upper and Lower Case Letters

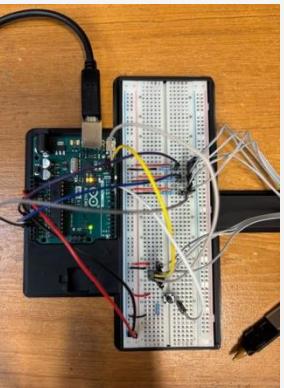
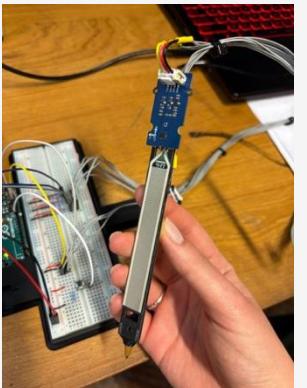
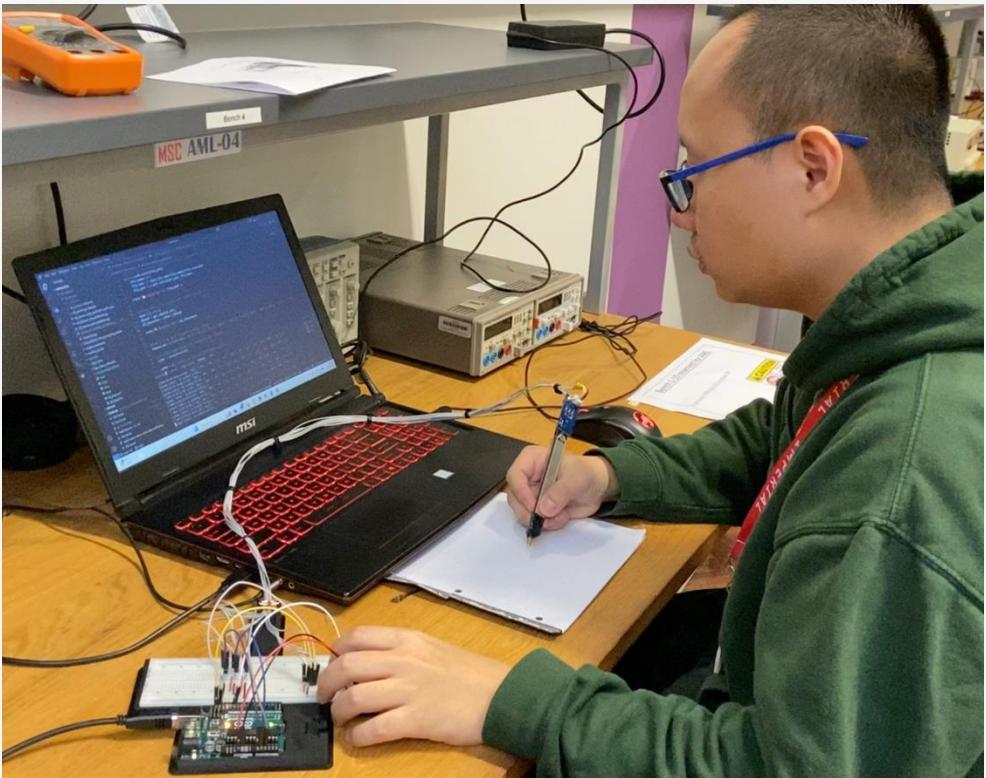
# Digits

# Maths Symbols

## Special Characters

## Words & Phrases

# Data Collection



```
IMU_to_CVS.py x
1.00
Users > tunakisaga > Desktop > IMU_to_CVS.py > ...
20 noise = ['noise']
21 english_alphabet_capital = ['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z']
22 english_alphabet_lower = ['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z']
23 numbers = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
24
25 all_characters = noise + english_alphabet_capital + english_alphabet_lower + numbers

PROBLEMS 7 OUTPUT TERMINAL COMMENTS
660, -435, 616, -1, -2, 1, -13, -12, 8
Current letter: F, writing to file...
Recording stopped.
654, -327, 594, 13, 2, 14, -13, -16, 10
Current letter: G, writing to file...
714, -348, 594, -4, -22, -5, -12, -16, 10
Current letter: G, writing to file...
761, -361, 552, -20, -14, -7, -12, -17, 11
Current letter: G, writing to file...
666, -242, 550, -26, -38, -16, -13, -18, 13
Current letter: G, writing to file...

Python - tunakisaga + x
```

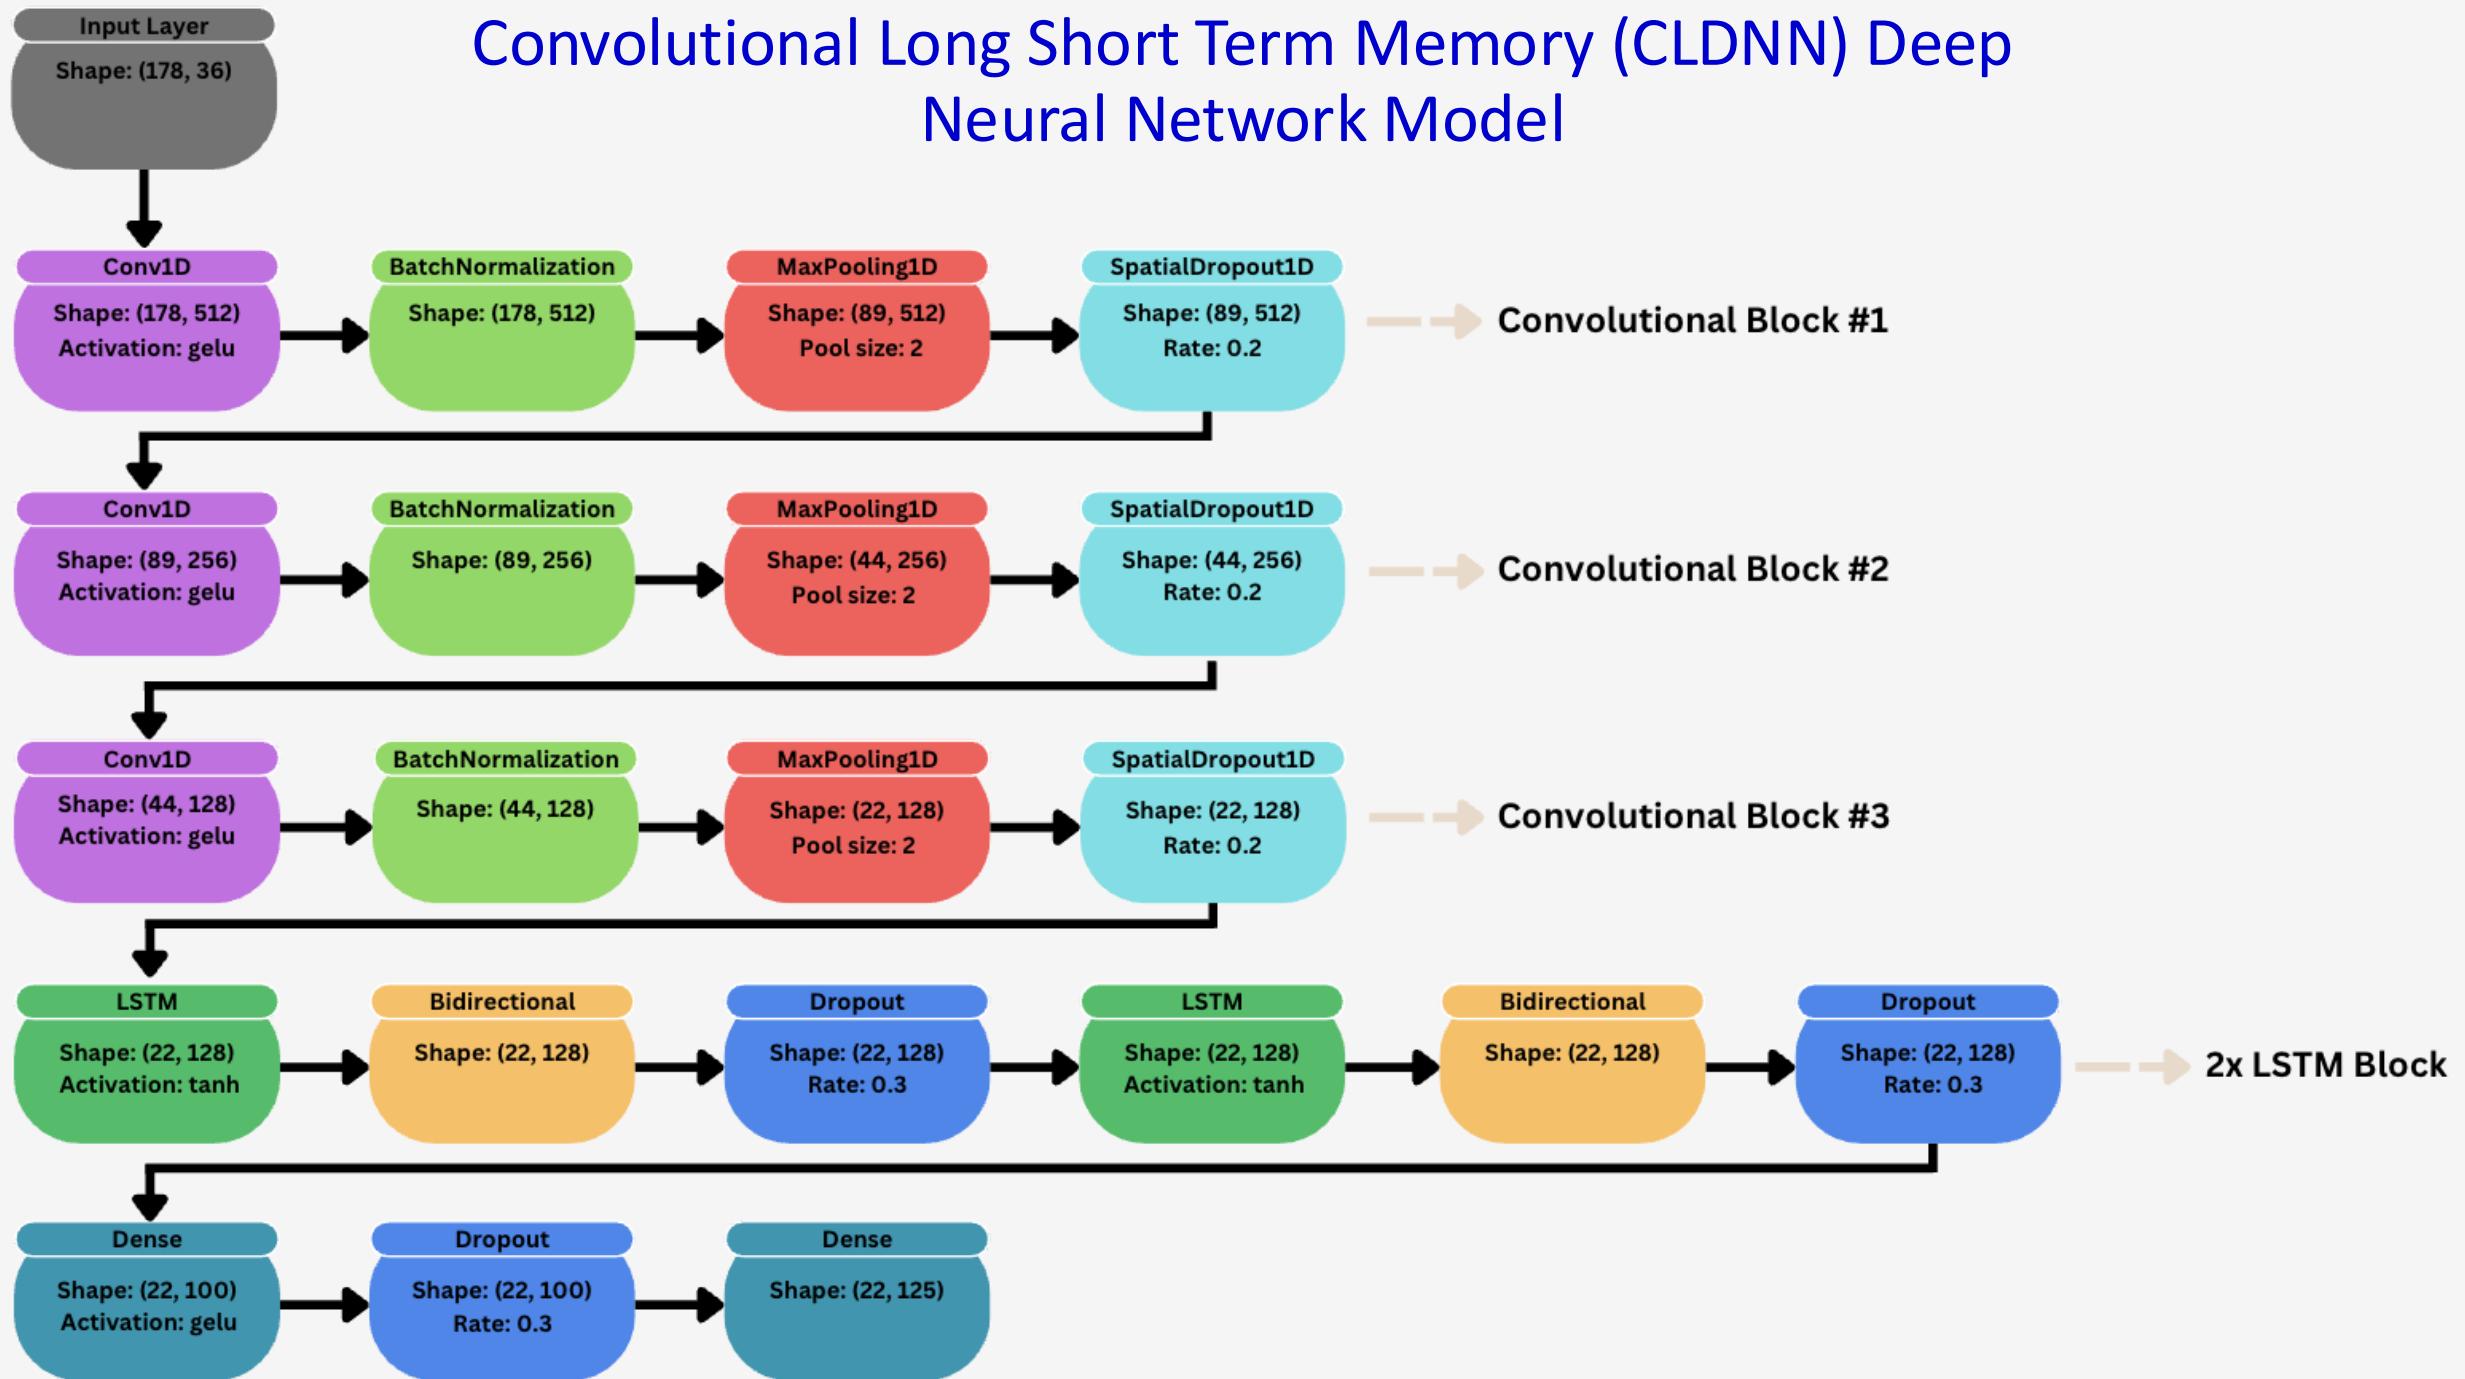
## Dataset Collection Protocol

- Use the final sensor-equipped pen and normal white paper on a flat, **horizontal surface**.
- Participants should sit comfortably at a table and maintain a consistent **pen orientation (IMU facing upwards)**.
- All English letters (uppercase and lowercase), digits, special characters and mathematical symbols should be included in the target label set.
- Each written character should be immediately **labeled** to ensure accurate ground truth.
- Each character should be written **naturally**, without strict size constraints.
- **Thresholding** should be applied to trim unnecessary sensor data at the start/end of strokes.
- The sensor data, timestamps, and labels should be **stored** in a structured format (e.g., CSV).

# Data Preprocessing

1. Remove non-writing data	2. Scale the data	3. Create sequences	4. Map characters
<p>Filter out non-writing sequences of data to ensure clean input signals for the CLDNN model. Non-writing sequences were determined using the force readings.</p> <pre># Filter out noise filtered_data = combined_data[     combined_data["Letter"] != "noise"]</pre>	<p>Use StandardScaler to normalize IMU and force sensor readings to zero mean and unit variance, ensuring stable training dynamics in the CLDNN model.</p> <pre># Normalize sensor data (Z-score normalization) scaler = StandardScaler() filtered_data[sensor_columns] = scaler.fit_transform(     filtered_data[sensor_columns])</pre>	<p>Convert raw time-series data into sequences of sensor data and label the sequences.</p> <pre>for _, row in data.iterrows():     if row["Letter"] != current_label:         # Save the current sequence if it's not empty         if current_sequence:             sequences.append(np.array(current_sequence))             labels.append(current_label)         # Start a new sequence         current_label = row["Letter"]         current_sequence = []     # Append sensor values to the current sequence     current_sequence.append(row[sensor_columns].values)</pre>	<p>Utilize a StringLookup layer to convert character labels into integer-encoded sequences for the CTC loss function.</p> <pre># Define a blank token blank_token = 'BLANK'  # Update the StringLookup layers char_to_num = layers.StringLookup(     vocabulary=list(characters),     mask_token=None,     oov_token=blank_token )</pre>
5. Introduce the blank token	6. Pad labels and sequences	7. Filter the sequences	8. Split the data
<p>Insert a dedicated blank index required by CTC to handle spacing and transitions between characters. Use 'BLANK' as the blank token.</p> <pre># Define a blank token blank_token = 'BLANK'</pre>	<p>Apply padding to sensor data sequences and corresponding labels to ensure uniform input size before feeding into the CLDNN model.</p> <pre># Pad sequences to the same length padded_sequences = pad_sequences(     filtered_sequences,     maxlen=max_length,     padding='post',     dtype='float32',     value=0)</pre>	<p>Remove sequences that are shorter than the minimum length to maintain dataset consistency.</p> <pre># Filter out very short sequences filtered_sequences = [seq for seq in sequences                      if seq.shape[0] &gt;= min_length] filtered_labels = [label for seq, label in zip(     sequences, padded_labels) if                      seq.shape[0] &gt;= min_length]</pre>	<p>Partition the processed dataset into train, test and validation sets using the designated function. The dataset is split into 70% training, 20% validation, and 10% test sets..</p> <pre># Apply the fixed split_data function X_train, X_test, y_train, y_test = split_data(     padded_sequences_tensor,     filtered_labels)</pre>

# Convolutional Long Short Term Memory (CLDNN) Deep Neural Network Model



# Model

## Pretraining



# STABILØ

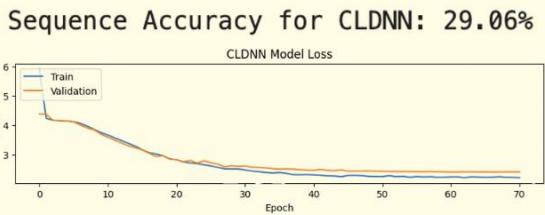
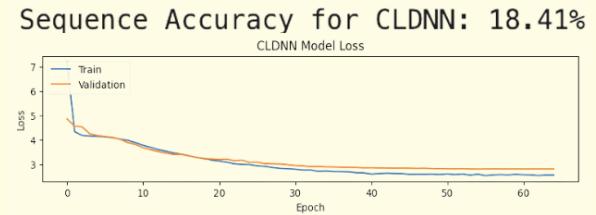
## The OnHW Datasets

- Front accelerometer (STM LSM6DSL)
- Gyroscope (STM LSM6DSL)
- Rear accelerometer (Freescale MMA8451Q)
- Magnetometer (ALPS HSCDTD008A)
- Force sensor (ALPS HSFPAR003A)

- 1 Train the model with the OnHW dataset, only including the IMU sensor data [6]
- 2 When training the model again with our data, transfer the weights from the first model to the IMU-related layers of the new model

Very Small Set of Data

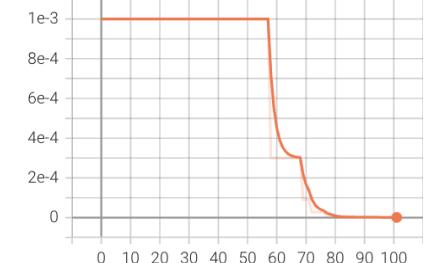
Without Pretraining



## Learning rate scheduling

Adjusts the learning rate dynamically during training (using ReduceLROnPlateau) to optimize convergence and prevent overshooting.

epoch\_learning\_rate  
tag: epoch\_learning\_rate



## Early stopper

Monitors validation loss and halts training if no improvement is observed for a set number of epochs, preventing overfitting.

```
Epoch 81/500
32/32    3s 89ms/step - loss: 5.5545 - val_loss: 8.8750
Epoch 82/500
32/32    5s 91ms/step - loss: 5.4869 - val_loss: 8.8620
Epoch 83/500
32/32    3s 106ms/step - loss: 5.5524 - val_loss: 8.8959
Epoch 84/500
32/32    5s 92ms/step - loss: 5.4954 - val_loss: 8.8613
Epoch 85/500
32/32    5s 91ms/step - loss: 5.5793 - val_loss: 8.8739
Epoch 86/500
32/32    5s 91ms/step - loss: 5.4945 - val_loss: 8.8721
Epoch 87/500
32/32    3s 89ms/step - loss: 5.6217 - val_loss: 8.8716
Epoch 88/500
32/32    3s 89ms/step - loss: 5.5147 - val_loss: 8.8722
Epoch 89/500
32/32    4s 122ms/step - loss: 5.4756 - val_loss: 8.8701
Epoch 90/500
32/32    4s 93ms/step - loss: 5.4861 - val_loss: 8.8711
Epoch 91/500
32/32    3s 98ms/step - loss: 5.5345 - val_loss: 8.8709
Epoch 92/500
32/32    6s 113ms/step - loss: 5.5207 - val_loss: 8.8712
Epoch 93/500
32/32    3s 91ms/step - loss: 5.5483 - val_loss: 8.8708
```

## Custom CTC loss

```
... Title
1 def ctc_loss(y_true, y_pred):
2
3     # Transpose y_pred
4     y_pred = tf.transpose(y_pred, [1, 0, 2])
5
6     # Ensure y_true is of type int32
7     y_true = tf.cast(
8         y_true, tf.int32)
9
10    # Calculate input length (logit length)
11    # and label length
12    logit_length = tf.fill(
13        [tf.shape(y_pred)[1]],
14        tf.shape(y_pred)[0]) # shape:
15    batch_size_length = tf.reduce_sum(
16        tf.cast(
17            tf.not_equal(y_true, 0), tf.int32
18        ), axis=1) # shape: (batch_size,)
19
20    # Compute the CTC loss
21    # using tf.nn.ctc_loss
22    loss = tf.nn.ctc_loss(
23        labels=y_true,
24        logits=y_pred,
25        label_length=label_length,
26        logit_length=logit_length,
27        logits_time_major=True,
28        blank_index=0
29    )
30
31    return loss
```

# Model Optimisation

1

**Before**

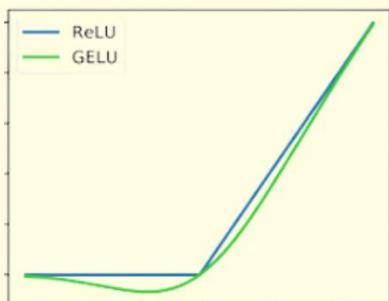
Conv1D and Dense layers:

```
x = layers.Conv1D(256, 3,  
activation='relu',  
padding='same')(x)
```

**After**

Conv1D and Dense layers:

```
x = layers.Conv1D(256, 3,  
activation='gelu',  
padding='same')(x)
```

**Reason**

GELU provides smoother activation than ReLU and helps prevent dying neurons.

2

**Before**

Convolutional blocks:

```
x =  
layers.Dropout(0.3)(  
x)
```

**After**

Convolutional blocks:

```
x = layers.Spatial  
Dropout1D(0.2)(x)
```

**Reason**

Standard Dropout randomly removes individual units, while **SpatialDropout1D** drops entire feature maps along the channel axis. This helps preserve spatial dependencies in time-series data

3

**Before**

```
# Compile the model with CTC loss  
cldnn_model.compile(optimizer='adam',  
loss=ctc_loss)
```

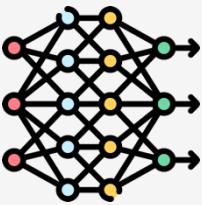
**After**

```
# Compile the model with CTC loss  
optimizer = tf.keras.optimizers.Adam(  
learning_rate=0.001, clipnorm=1.0)  
cldnn_model.compile(optimizer=optimizer,  
loss=ctc_loss)
```

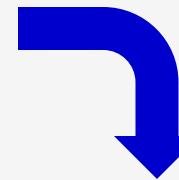
**Reason**

Gradient clipping (clipnorm=1.0) stabilizes training by preventing large updates that could disrupt convergence. This enhances model robustness and prevents exploding gradients.

# Model Optimisation



CLDNN model [5]



1

Introduce derivatives

2

Replace relu with gelu

3

Set clipnorm to 1

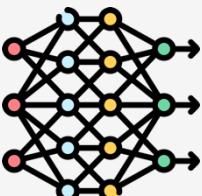
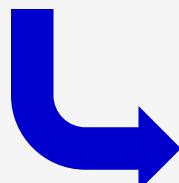
4

Add learning rate scheduler

5

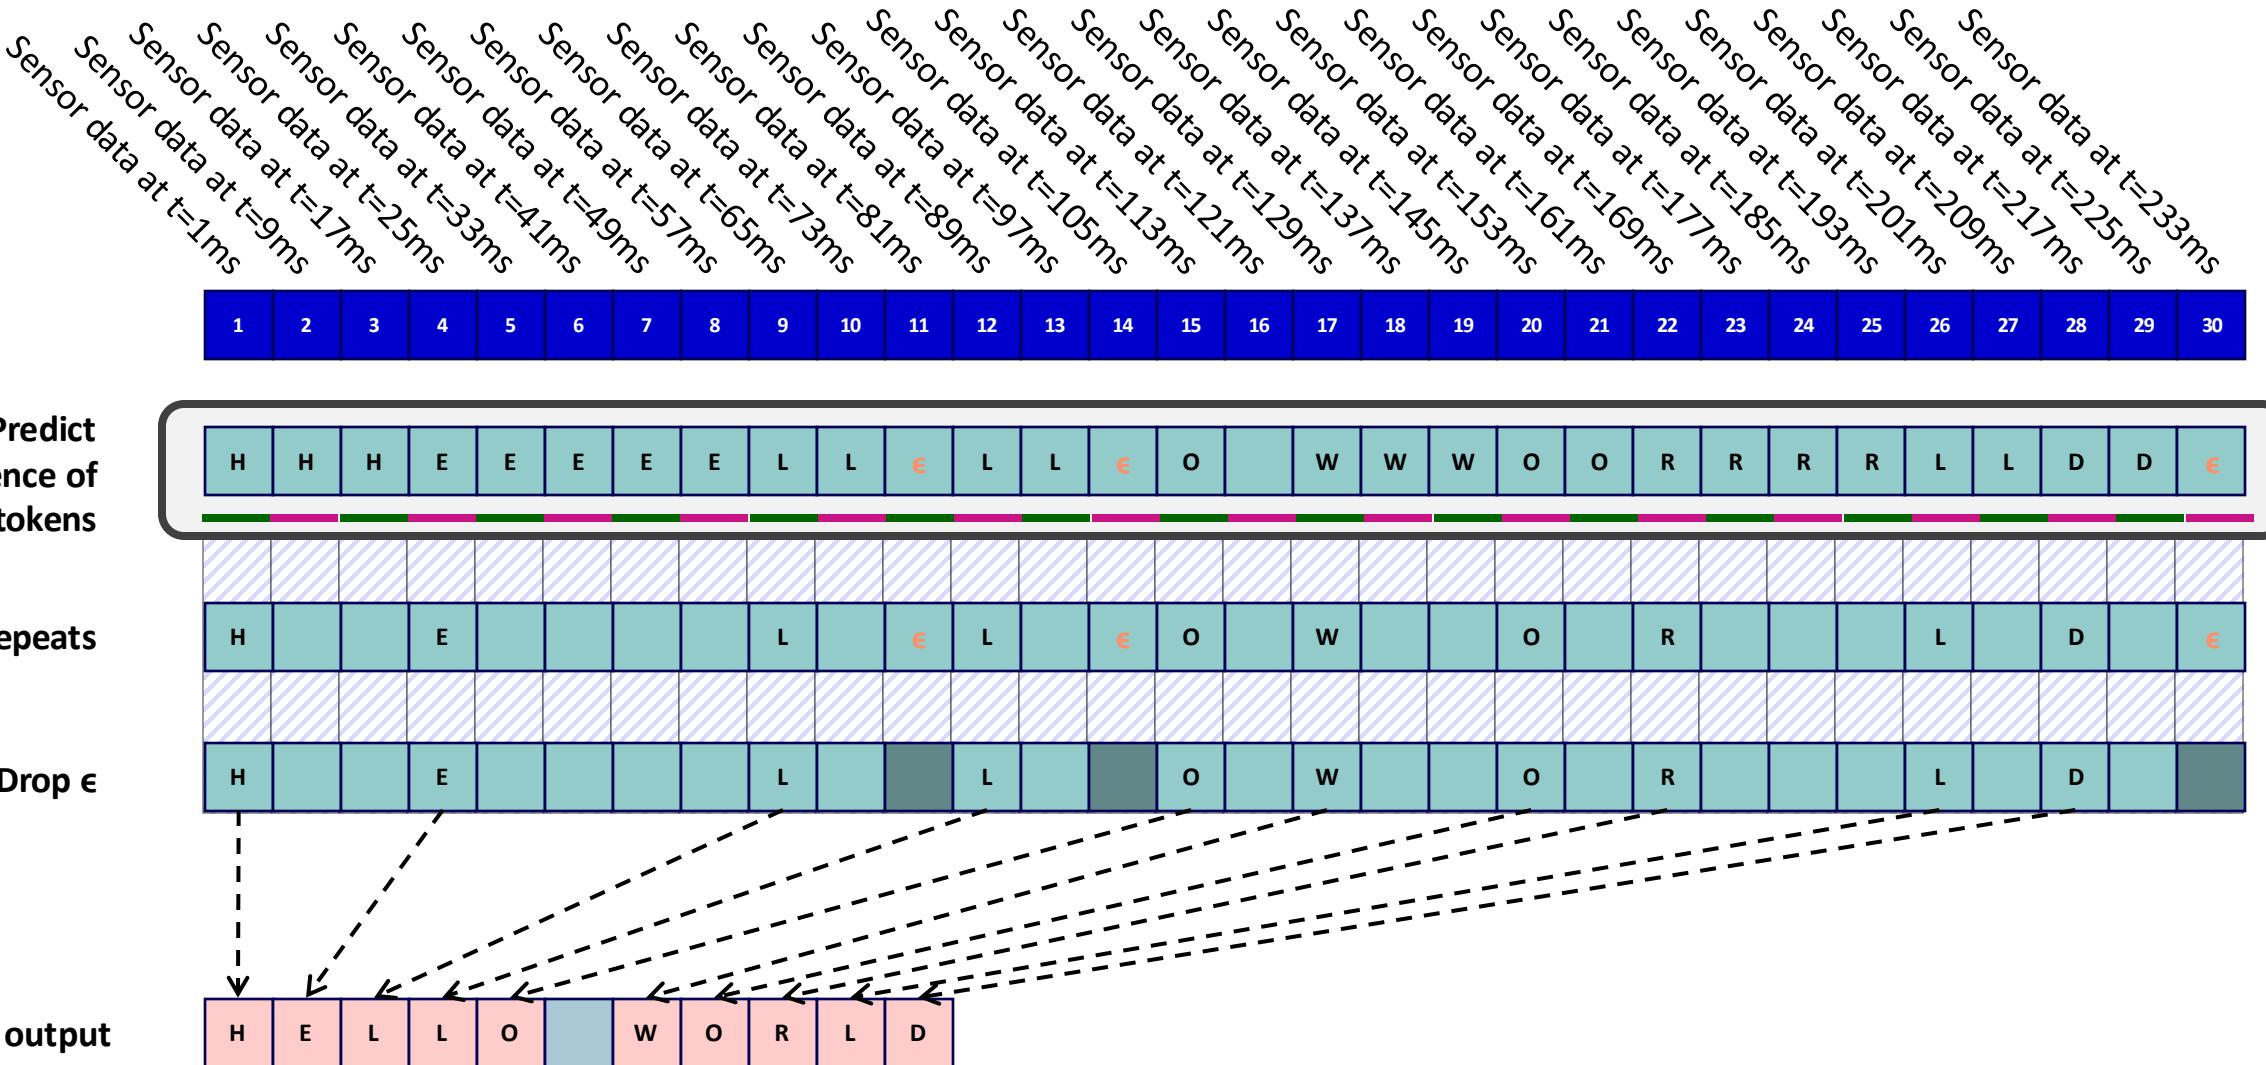
Replace dropout with spatial dropout

Change	Accuracy Before the Change (%)	Accuracy After the Change (%)	Step Improvement (%)	Cumulative Improvement (%)
1	5.63	6.78	1.15	1.15
2	6.78	12.93	6.15	7.3
3	12.93	20.35	7.42	14.72
4	20.35	58.08	37.73	52.45
5	58.08	63.15	5.07	57.52

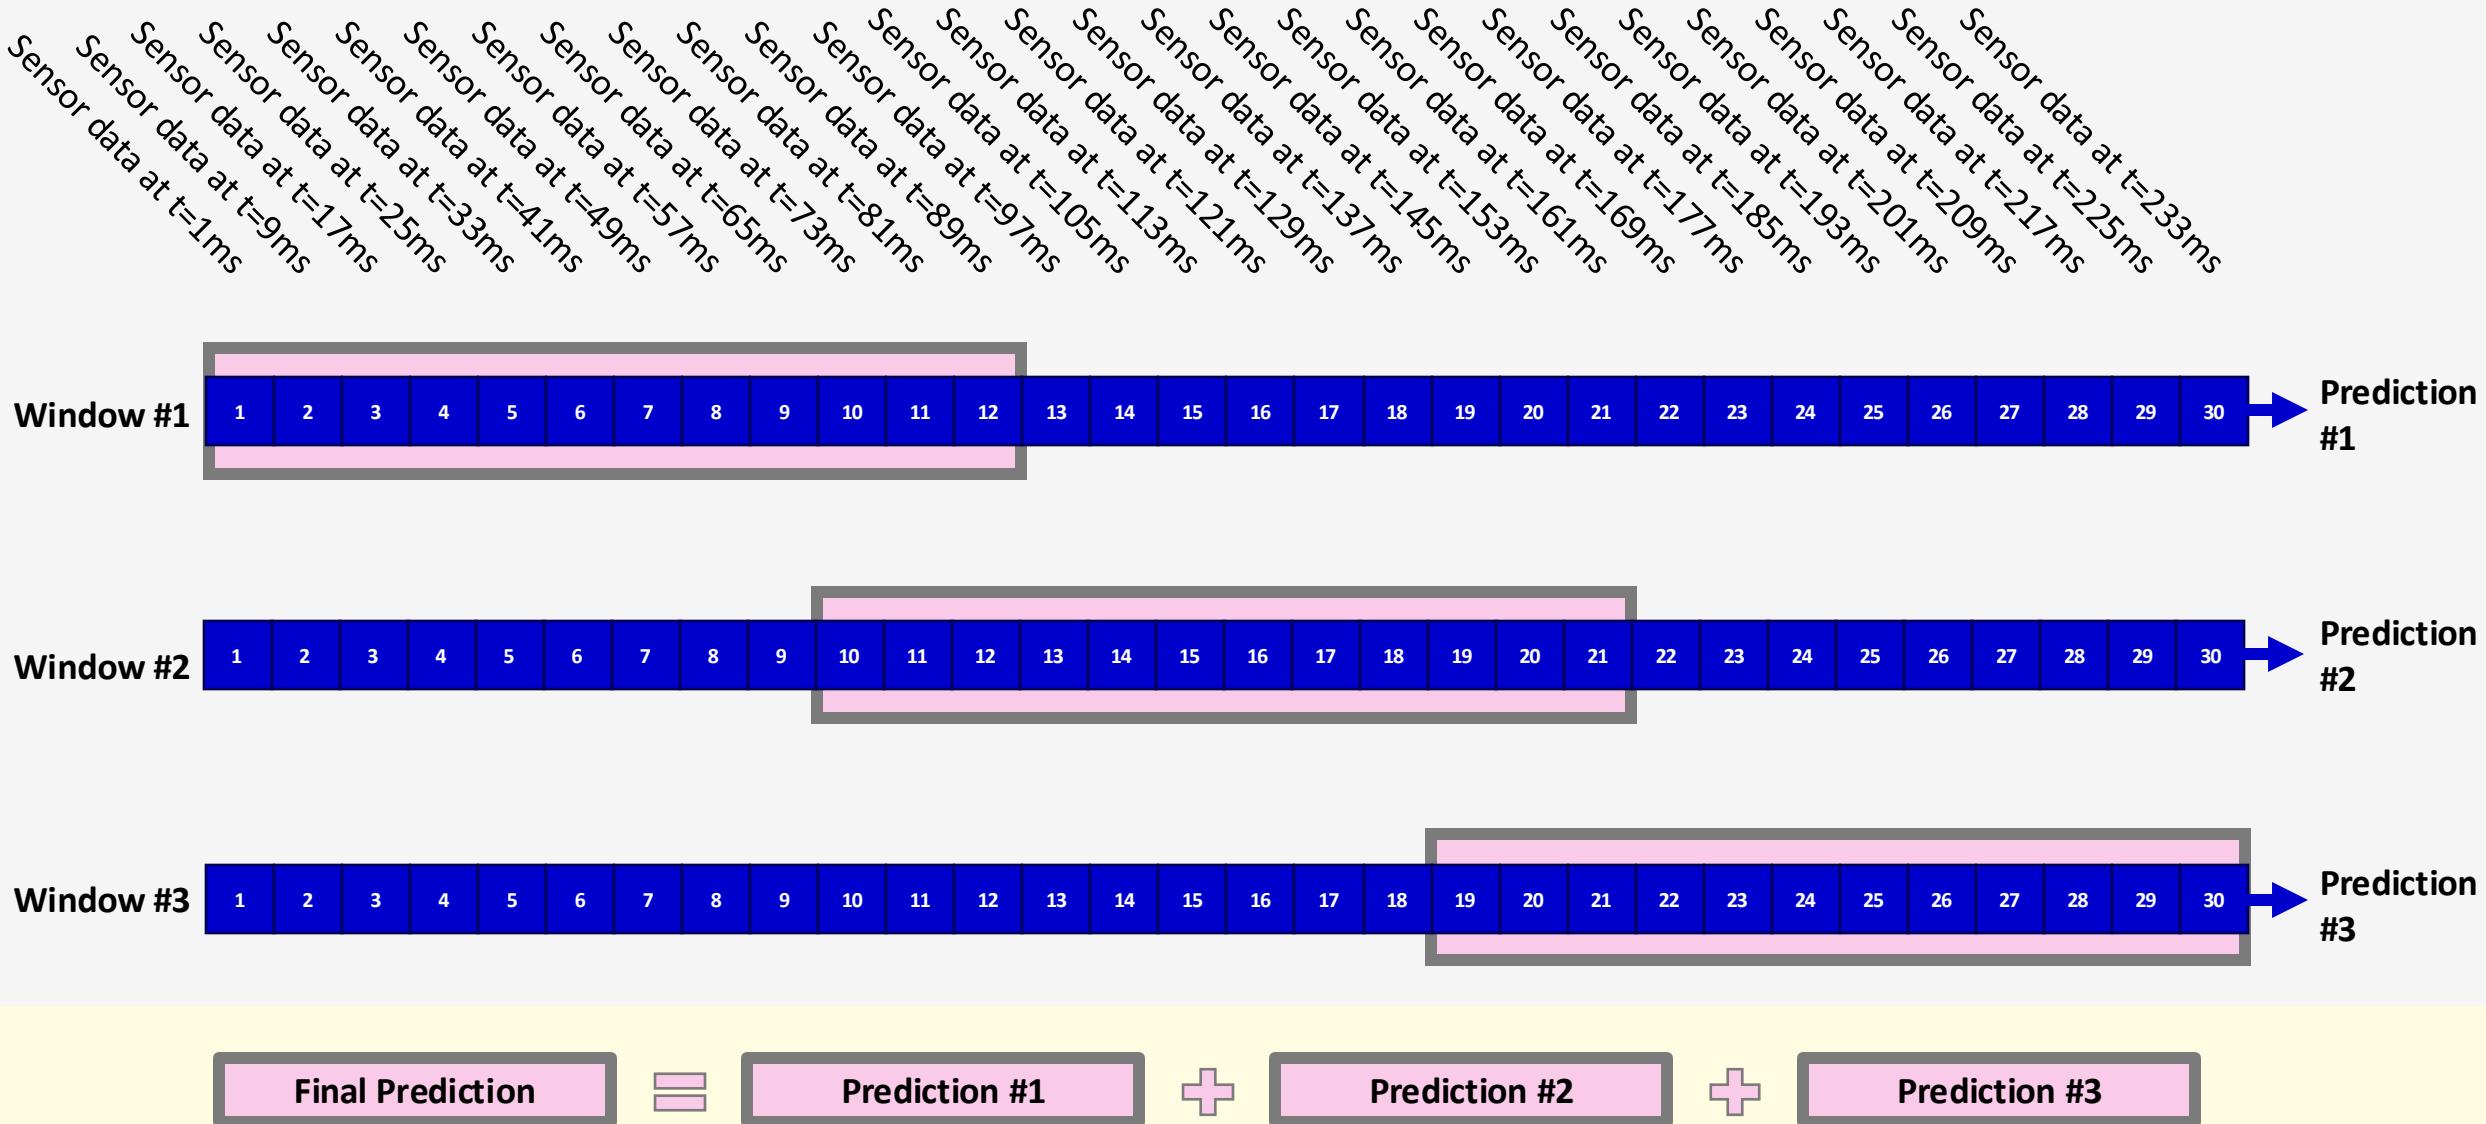


Optimised CLDNN model with improved accuracy

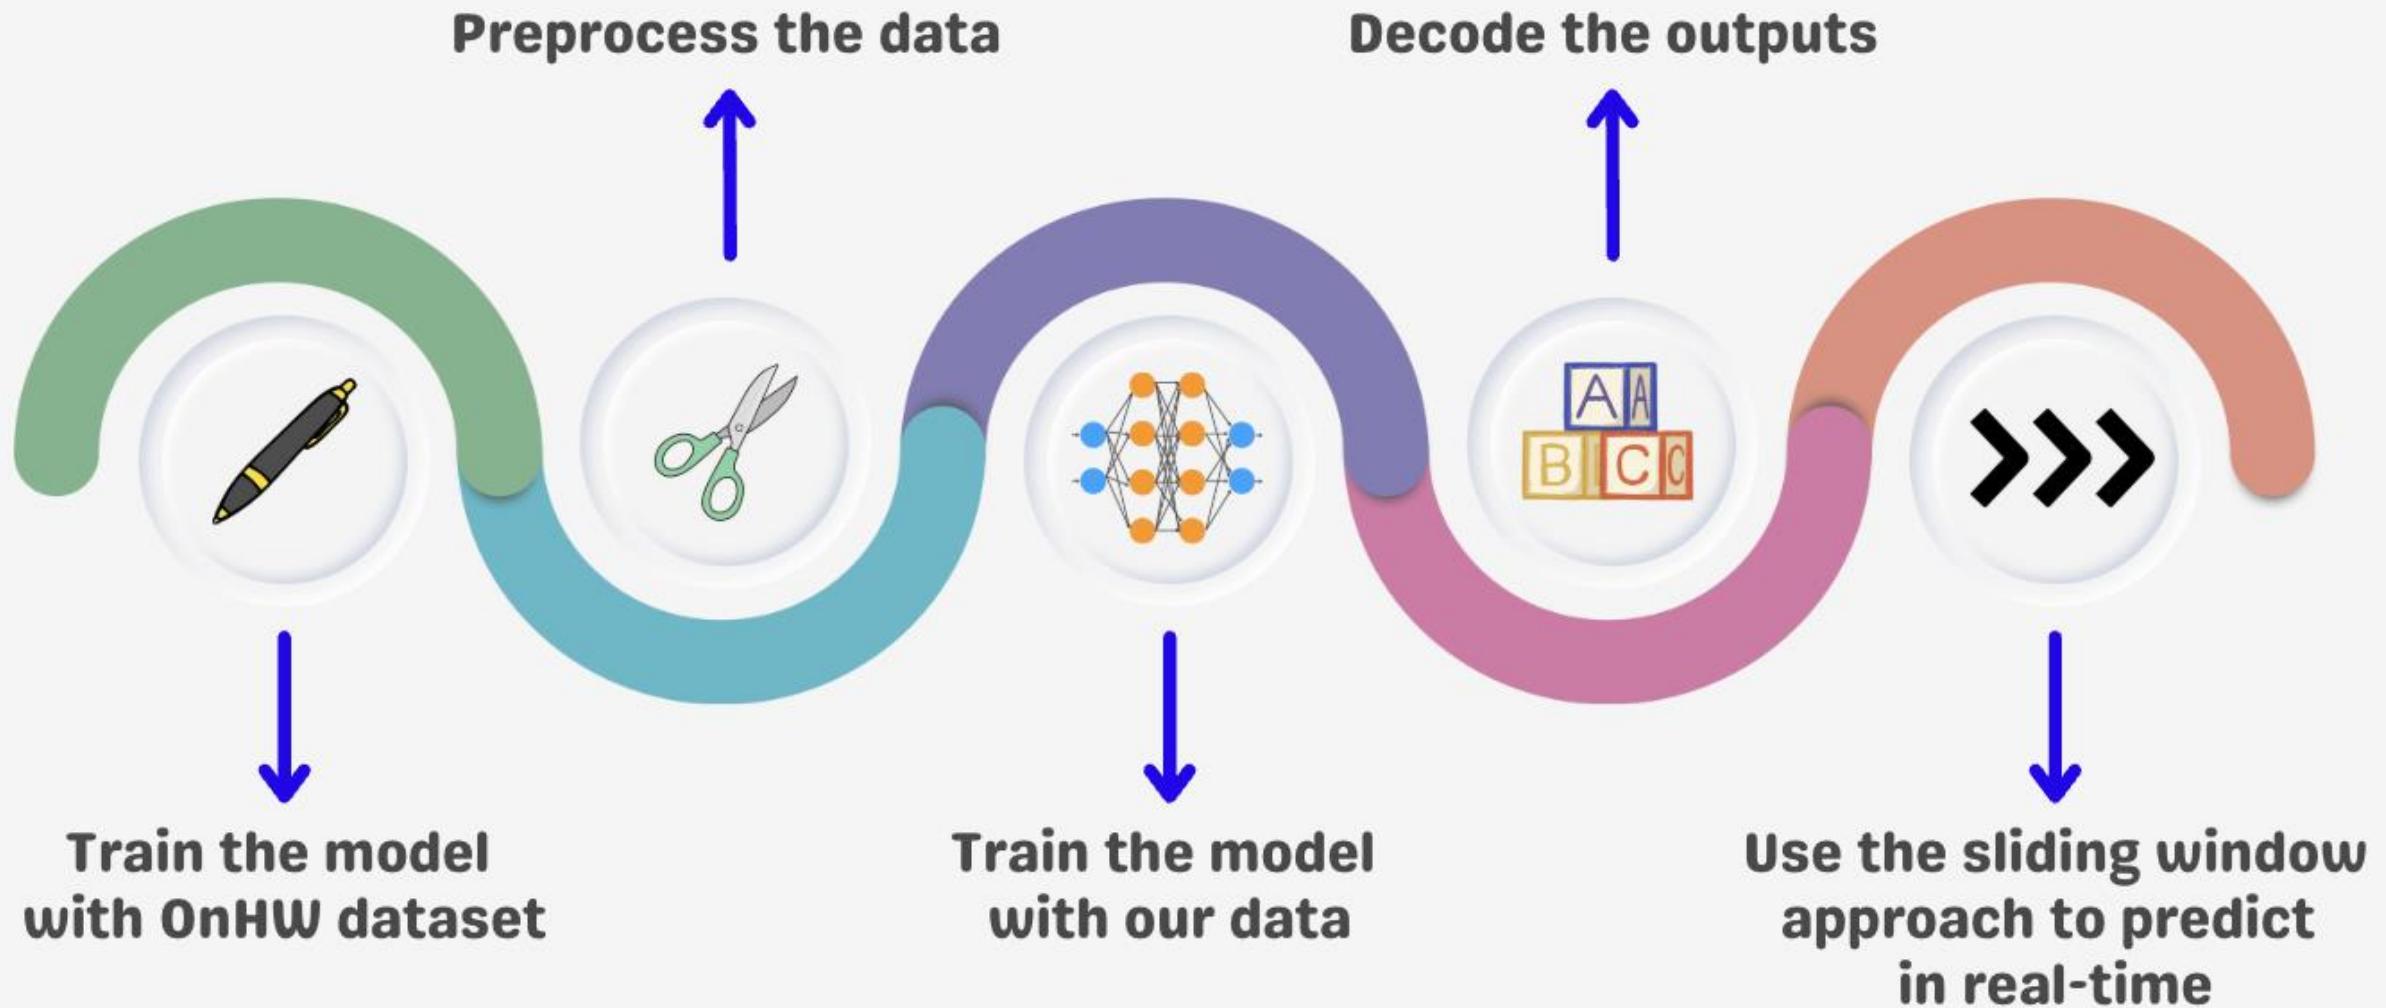
# CTC Loss



# Sliding Window Approach



# End-to-end Machine Learning Pipeline



# Experimental Results

## Model Performances, Loss Plots

# Discussion on Why Some Models Outperform Others

```
# Sort models (highest sequence accuracy first)
sorted_models = sorted(results.items(), key=lambda x: x[1], reverse=True)

# Print
for model_name, accuracy in sorted_models:
    print(f"{model_name}: Sequence Accuracy = {accuracy * 100:.2f}%")
```

CLDNN: Sequence Accuracy = 23.04%

TCN-BiLSTM-Attention Hybrid: Sequence Accuracy = 21.53%

CNN-BiLSTM-Attention: Sequence Accuracy = 7.64%

Attention-Based Seq2Seq Model: Sequence Accuracy = 5.81%

InceptionTime: Sequence Accuracy = 4.74%

TCN: Sequence Accuracy = 3.88%

CNN: Sequence Accuracy = 1.29%

BiLSTM: Sequence Accuracy = 0.86%

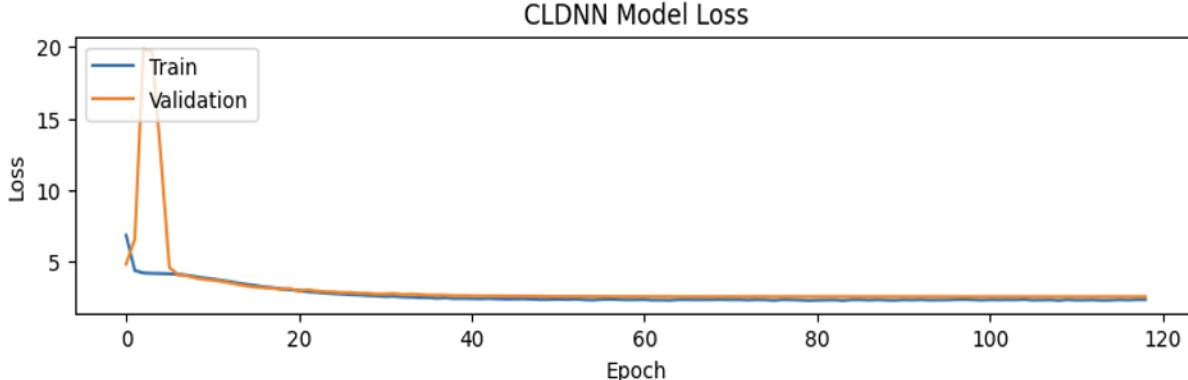
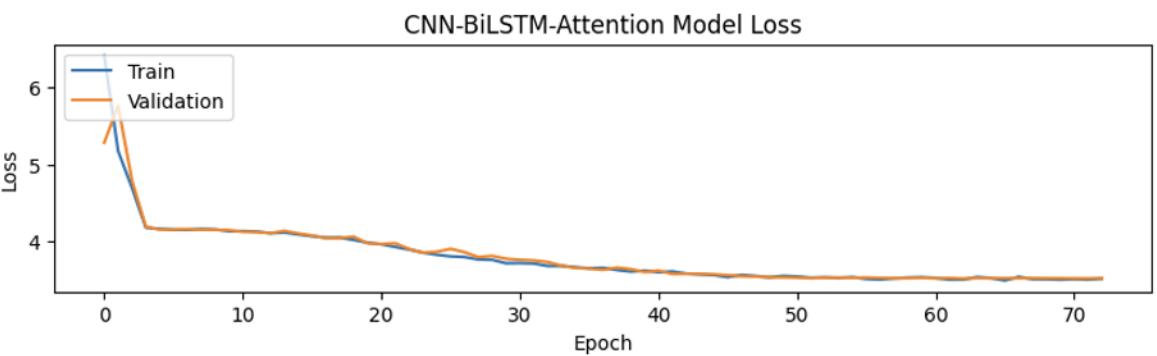
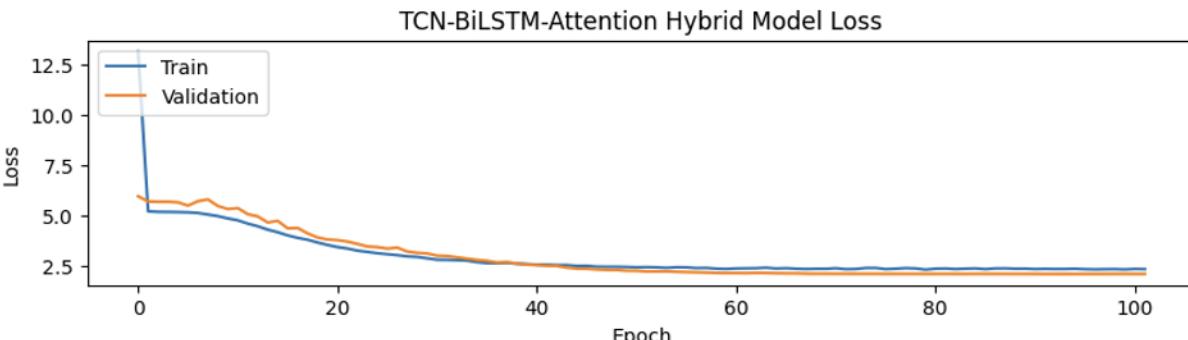
HMM-RNN: Sequence Accuracy = 0.86%

BLSTM-CTC: Sequence Accuracy = 0.86%

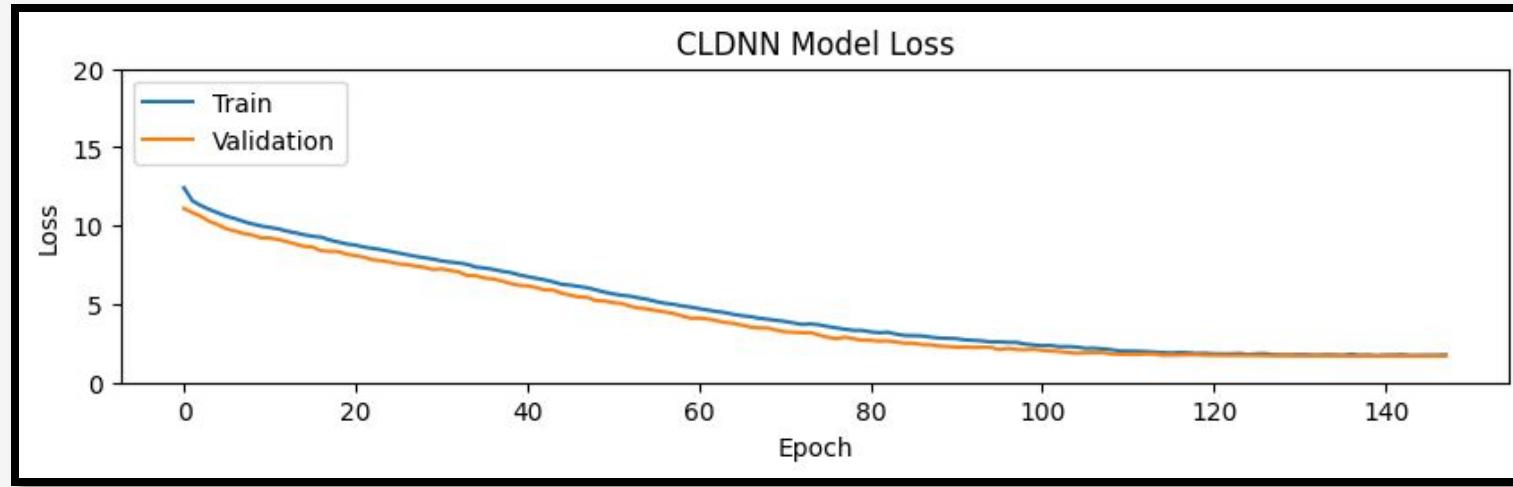
Transformer: Sequence Accuracy = 0.00%

TCN-BiLSTM-Transformer Hybrid: Sequence Accuracy = 0.00%

Model	Acc. (%)	Key Strengths
CLDNN	23.04	Strong <b>feature extraction</b> (CNN) + <b>temporal dependencies</b> , past and future contexts (LSTM)
TCN-BiLSTM-Attention Hybrid	21.53	<b>Long-term dependencies</b> without losing information about the <b>temporal order</b> (TCN), preceding and succeeding sensor readings (BiLSTM), <b>critical movements</b> , such as <b>sharp turns</b> in the pen trajectory or pressure changes (attention)
CNN-BiLSTM-Attention	7.64	Effective hybrid approach, but weaker attention



# Training and Validation Losses



Decoding Method	Sequence Accuracy (%)
Greedy	61.97
Beam	68.13

Sample True vs. Predicted Comparisons:

True: hbd | Predicted: hd  
True: l | Predicted: l  
True: B | Predicted: B  
True: BLANK | Predicted: BLANK  
True: f | Predicted: x  
True: dkdog ms | Predicted: dnog m  
True: XdkknvZSqdd | Predicted: XdknvZSd  
True: c | Predicted: T  
True: z | Predicted: N  
True: bgdqqx | Predicted: bgdqq  
True: | | Predicted:  
True: C s ZRbhmbd | Predicted: C ZRbhmb  
True: kdlnm | Predicted: kdlnm  
True: t | Predicted: t  
True: S | Predicted: S  
True: 4 | Predicted: 4  
True: G | Predicted: L  
True: uhnkham | Predicted: mhfkhan  
True: 4 | Predicted: 4  
True: G | Predicted: G

# Workplan

## Milestones and Challenges

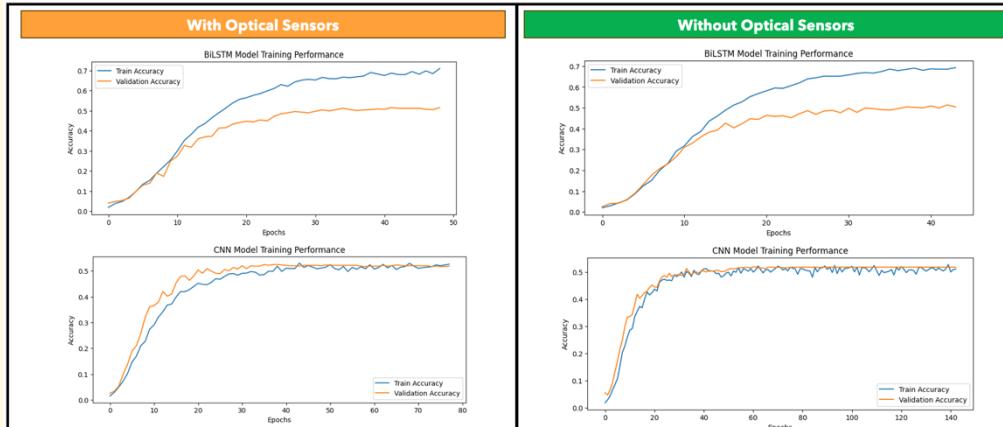
# Project Timeline

## Gantt Chart

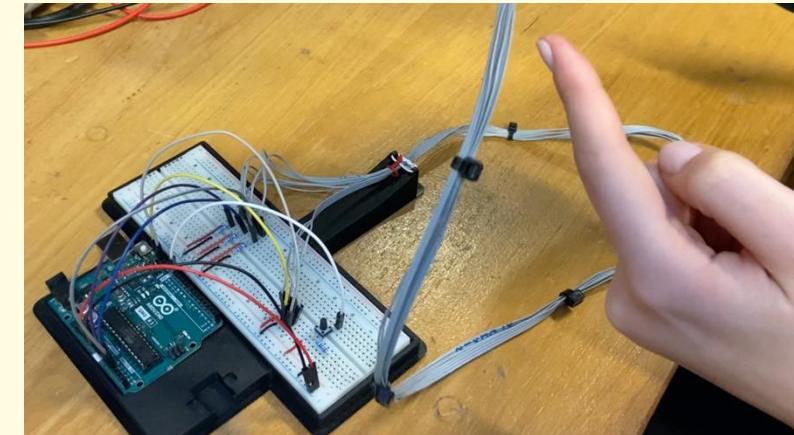


# Key Challenges and Decisions

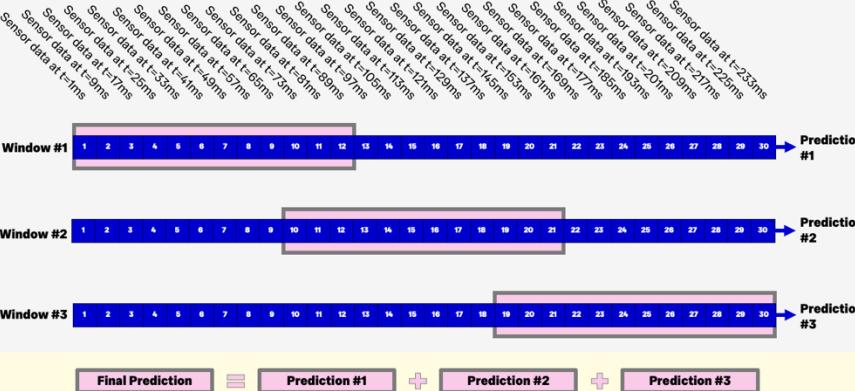
## Removing the optical sensor



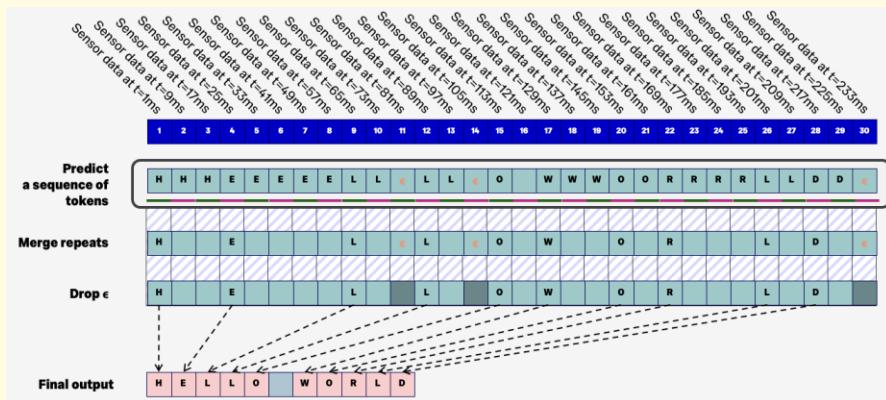
## Long wiring



## Real time decoding



## CTC loss being different from what we've seen so far



# Conclusion

## Summary, Evaluation, Future Work

# Summary & Evaluation

To Summarize What We Have Done:

Able to implement a CLDNN model for our project

Able to collect handwriting dataset using the hardware

Able to have a working final project overall

Able to create a working smart pen hardware

Able to optimize the model that used our dataset

# Potential Future Work

**Extended Dataset Collection**  
Increase the dataset size by collecting more handwriting samples from diverse participants, including different writing styles, speeds, and orientations.

**Context-Aware Recognition**  
Implement a language model (e.g., LLM or transformer-based post-processing) to improve recognition by leveraging word and sentence context.

**Handwriting-to-Editable-Document**

Expand beyond LaTeX to support real-time handwriting-to-Word or Markdown conversion for broader usability.

**Few-Shot Learning for Personalization**  
Train models that adapt to individual handwriting styles with minimal additional data using meta-learning techniques.

**Erasure and Editing Improvements**  
Enhance dynamic erasure detection and correction, allowing seamless real-time modifications of handwritten text.

**Multi-Language Support**

Extend the system to support multiple languages and scripts, adapting character recognition models accordingly.

# References

## Papers and Other Resources

## References

- [1] A. Author, B. Author, and C. Author, "A Self-Attention-Based Deep Architecture for Online Handwriting Recognition," *Proc. Conf. on Pattern Recognition*, 2024.
- [2] X. Zhang, Y. Liu, and Z. Wang, "Representing Online Handwriting for Recognition in Large Vision-Language Models," *IEEE Trans. Pattern Anal. Mach. Intell.*, 2024.
- [3] M. Wehbi, T. Hamann, J. Barth, P. Kaempf, D. Zanca, and B. Eskofier, "Benchmarking Online Sequence-to-Sequence and Character-Based Handwriting Recognition from IMU-Enhanced Pens," in *Proc. ICDAR*, 2022.
- [4] P. Smith and R. Lee, "Improving Accuracy and Explainability of Online Handwriting Recognition," *Neural Computation*, vol. 34, no. 3, pp. 1234–1250, 2022.
- [5] M. Wehbi, T. Hamann, J. Barth, P. Kaempf, D. Zanca, and B. Eskofier, "Towards an IMU-Based Pen Online Handwriting Recognizer," in *Document Analysis and Recognition – ICDAR 2021*, J. Lladós, D. Lopresti, and S. Uchida, Eds. Cham: Springer, 2021, vol. 12823, Lecture Notes in Computer Science. doi: 10.1007/978-3-030-86334-0\_19.
- [6] F. Ott, M. Wehbi, T. Hamann, J. Barth, B. Eskofier, and C. Mutschler, "The OnHW Dataset: Online Handwriting Recognition from IMU-Enhanced Ballpoint Pens with Machine Learning," *Proc. ACM Interact. Mob. Wearable Ubiquitous Technol. (IMWUT)*, vol. 4, no. 3, article 92, pp. 1–20, Sep. 2020, doi: 10.1145/3411842.

## References

- [7] L. Kim, H. Park, and J. Choi, “Deep-Learning-Based Character Recognition from Handwriting Motion Data Captured Using IMU and Force Sensors,” *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 33, no. 9, pp. 5231–5243, 2022.
- [8] A. Gupta and S. Mehta, “Motion-Based Handwriting Recognition and Word Reconstruction,” *Pattern Recognition Letters*, vol. 151, pp. 12–21, 2021.
- [9] Voidful, “Understanding CTC loss for speech recognition,” *Medium*, [Online]. Available: <https://voidful.medium.com/understanding-ctc-loss-for-speech-recognition-a16a3ef4da92>. [Accessed: Mar. 20, 2025].
- [10] Wikipedia, “Rectifier (neural networks),” *Wikipedia*, [Online]. Available: [https://en.wikipedia.org/wiki/Rectifier %28neural\\_networks%29](https://en.wikipedia.org/wiki/Rectifier_%28neural_networks%29). [Accessed: Mar. 20, 2025].
- [11] Distill, “CTC,” *Distill.pub*, [Online]. Available: <https://distill.pub/2017/ctc/>. [Accessed: Mar. 20, 2025].

# IMPERIAL

## Thank you